

הארכה בשל מילואים בצו 8 של שי - HW1

Shai Tayar 315034074

Gal Granot 315681593

Question 1: knowing the system

a) Cuda version: 12.5.40

```
$ nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2024 NVIDIA Corporation
Built on Wed Apr 17 19:19:55 PDT 2024
Cuda compilation tools, release 12.5, V12.5.40
Build cuda_12.5.r12.5/compiler.34177558_0
```

b) GPU name: NVIDIA GeForce RTX 2080

```
$ nvidia-smi
Mon Jun 24 06:27:14 2024
```

NVIDIA-SMI 550.90.07		Driver Version: 550.90.07		CUDA Version: 12.4	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC
Fan	Temp	Pwr:Usage/Cap		Memory-Usage	GPU-Util Compute M.
0	NVIDIA GeForce RTX 2080	Off	00000000:02:00.0	Off	N/A
30%	29C P8	10W / 250W	3M1B / 8192M1B	0%	Default N/A
1	NVIDIA GeForce RTX 2080	Off	00000000:03:00.0	Off	N/A
30%	28C P8	10W / 250W	3M1B / 8192M1B	0%	Default N/A
2	NVIDIA GeForce RTX 2080	Off	00000000:03:00.0	Off	N/A
32%	47C P2	77W / 250W	581M1B / 8192M1B	100%	Default N/A
3	NVIDIA GeForce RTX 2080	Off	00000000:04:00.0	Off	N/A
30%	27C P8	11W / 250W	3M1B / 8192M1B	0%	Default N/A

Processes:						
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage
2	N/A	N/A	383671	C	/home/u_208848499/homework1/ex1	136MiB
2	N/A	N/A	384817	C	/home/u_208848499/homework1/ex1	148MiB
2	N/A	N/A	444420	C	/home/u_208848499/homework1/ex1	148MiB
2	N/A	N/A	563586	C	/home/u_206902827/ex1	136MiB

c) After compiling and running deviceQuery:

```
Device 3: "NVIDIA GeForce RTX 2080 SUPER"
CUDA Driver Version / Runtime Version      12.4 / 12.5
CUDA Capability Major/Minor version number: 7.5
Total amount of global memory:              7967 MBytes (8354398208 bytes)
(048) Multiprocessors, (064) CUDA Cores/MP: 3072 CUDA Cores
GPU Max Clock rate:                        1815 Mhz (1.81 GHz)
Memory Clock rate:                         7751 Mhz
Memory Bus Width:                          256-bit
L2 Cache Size:                             4194304 bytes
Maximum Texture Dimension Size (x,y,z)      1D=(131072), 2D=(131072, 65536), 3D=(16384, 16384, 16384)
Maximum Layered 1D Texture Size, (num) layers 1D=(32768), 2048 layers
Maximum Layered 2D Texture Size, (num) layers 2D=(32768, 32768), 2048 layers
Total amount of constant memory:            65536 bytes
Total amount of shared memory per block:    49152 bytes
Total shared memory per multiprocessor:    65536 bytes
Total number of registers available per block: 65536
Warp size:                                  32
Maximum number of threads per multiprocessor: 1024
Maximum number of threads per block:        1024
Max dimension size of a thread block (x,y,z): (1024, 1024, 64)
Max dimension size of a grid size (x,y,z):  (2147483647, 65535, 65535)
Maximum memory pitch:                       2147483647 bytes
Texture alignment:                           512 bytes
Concurrent copy and kernel execution:       Yes with 3 copy engine(s)
Run time limit on kernels:                   No
Integrated GPU sharing Host Memory:          No
Support host page-locked memory mapping:    Yes
Alignment requirement for Surfaces:         Yes
Device has ECC support:                      Disabled
Device supports Unified Addressing (UVA):    Yes
Device supports Managed Memory:              Yes
Device supports Compute Preemption:          Yes
Supports Cooperative Kernel Launch:          Yes
Supports MultiDevice Co-op Kernel Launch:   Yes
Device PCI Domain ID / Bus ID / location ID: 0 / 132 / 0
Compute Mode:
< Default (multiple host threads can use ::cudaSetDevice() with device simultaneously) >
> Peer access from NVIDIA GeForce RTX 2080 SUPER (GPU0) -> NVIDIA GeForce RTX 2080 SUPER (GPU3): No
```

d) The number of multiprocessors is 48.

Question 2: implement device functions

- a. We implemented the `prefix_sum` device function, as we learned in class.

Question 3: Implement a task serial version

- a. We implemented the `process_image_kernel` function in our code.
- b. We used `AtomicAdd` in the function that computes the histogram.
In this function, all the threads in the thread block access the same array in shared memory and update it. In order to avoid race condition, in which two thread updates the same cell in the histogram simultaneously but the actual value increases in 1 (instead of 2), we use `AtomicAdd` instead of `histogram[index]++`.
`AtomicAdd` operation ensures that the increment operation is performed atomically, which means only one thread at a time can access the location and perform the addition.
- c. The global memory accesses are coalesced regardless of the number of the threads. The image's size is $N \times N$ and $N = k^2 (k \in \mathbb{N}), N \geq 64$. Therefore, the GPU access to the global memory is in full lines, since the threads access contiguous memory locations. Means that the memory accesses are coalesced, regardless the number of threads.
- d. The state needed: `uchar* all_in, uchar* all_out, uchar* map`.
- e. We allocated and released the necessary resources in the code.
- g. Each thread block can contain up to 1024 threads, as seen in question 1c. The number of threads should be a multiple of the tile width, but also not lower than `COLOR_RANGE` (256).

Moreover, the expression: $\frac{PIXELS_PER_TILE}{THREAD_NUM}$ should be an integer.

We decided to use 256 threads per thread block.

- h. The run time is 961.73 msec and the throughput is $\frac{1000}{578.04} \cdot 1000 = 1729.98 \frac{\text{images}}{\text{sec}}$

```
$ ./ex1
Number of devices: 4
Using device 3

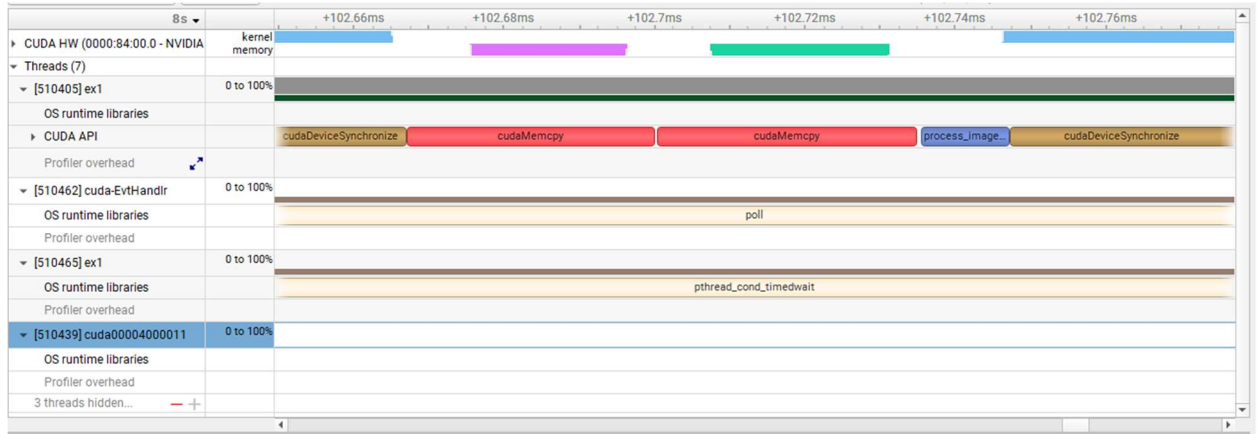
=== Randomizing images ===
total time 1597.933244 [msec]

=== CPU ===
total time 4970.291320 [msec]

=== GPU Task Serial ===
total time 578.048880 [msec] distance from baseline 0 (should be zero)

=== GPU Bulk ===
total time 49.816350 [msec] distance from baseline 0 (should be zero)
```

i. Execution diagram:



j.Memcpy from CPU to GPU, duration: 32.899 us

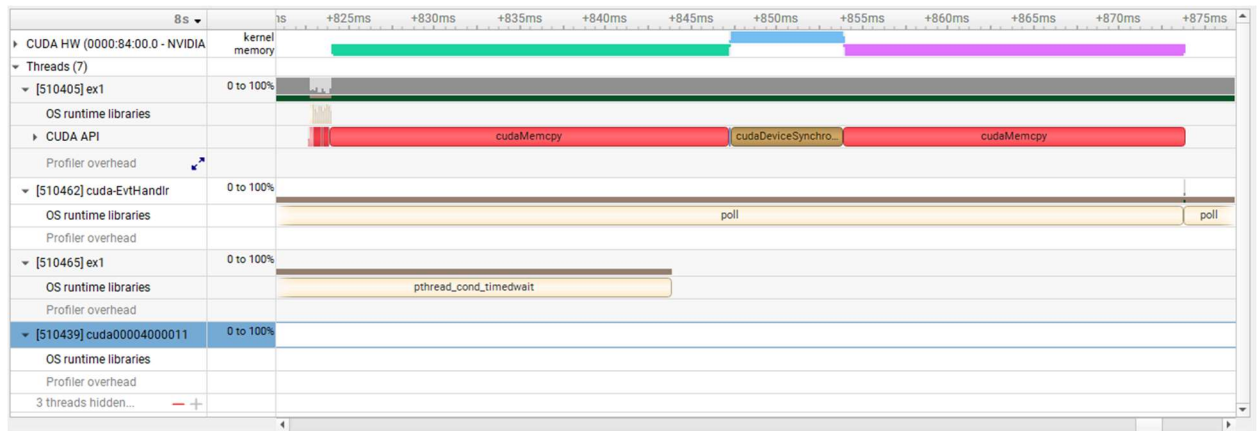
19	cudaMemcpy	8.10267s	32.899 µs	510405
20	cudaMemcpy	8.1027s	34.494 µs	510405

Question 4: Implement a bulk synchronous version

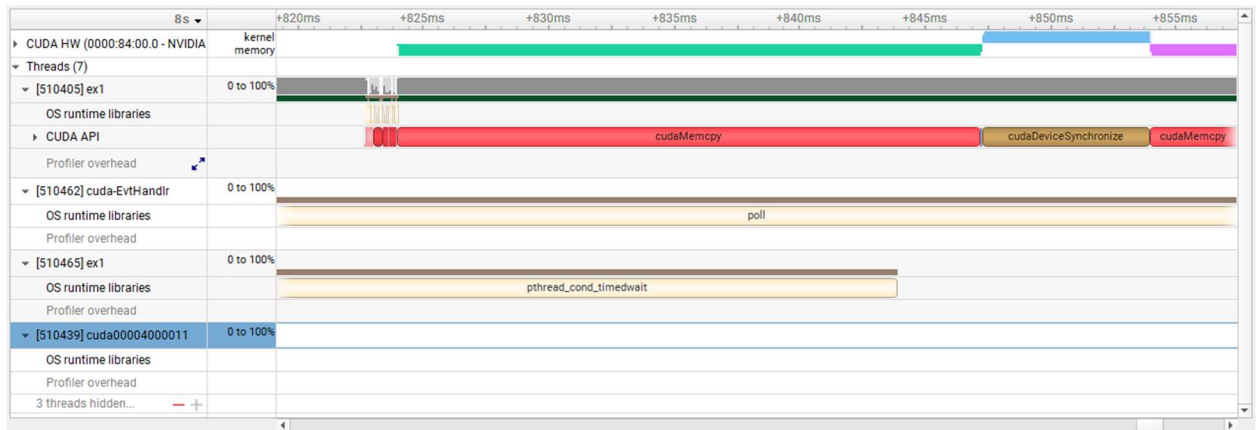
f. From image in Question 3h: Execution time: 49.816 msec.

$$\text{Speedup} = \frac{578.048}{49.816} = 11.603$$

g. Execution diagram:



h. CPU to GPU memcpc in the diagram:



Duration: 28.287 ms

$$\frac{Memcpy_{Bulk}}{Memcpy_{serial}} = \frac{28.287ms}{32.899 \mu s} = 859.81$$

The bulk version copied X1000 more images, which means the time does not grow linearly with the size of data being copied.